

ProgAlg - Résumé

Sylvain Julmy

9 juin 2016

Chapitre 1

Introduction

Concurrent computing : A form of computing in which programs are designed as a collections of interaction computational processes that may be executed in parallel.

Parallel computing : A form of computing in which many calculations are carried out simultaneously, operating on the principle that large problems can be divided into smaller ones. The most popular form is the multi-core processors.

Distributed computing : Any computing that involves multiple computers remote from each other that each have a role in a computation problem or information processing. Motivation : High performance (HPC) to avoid processing issue (compute faster) and memory issue (compute larger problems).

There is different forms of parallel computing :

- Bit-level
- Instruction level
- Data parallelism
- Task parallelism

There is 3 ways to get performance :

- More powerful processing unit (CPU)
- Better algorithms
- Parallelize / distribute

N computers performing the task T would compute the task faster than with 1 computer but never N times faster.

Flops : number of floating point operation over time. $1Flops = 1$ floating point operation in 1 second.

MIPS : million of instructions per second. $1MIPS = 1$ million of instructions in one second

Peak performance : the maximum computing performance of a computer.

Sustained performance : the observed performance on a benchmark program (usually LINPACK).

Due to the rapid increasing of the numbers of connected processors it rapidly appeared that purely "shared memory" supercomputer is not an option (memory bottleneck).

- Processing unit : Smallest sequential processing element, a core.
- SMP : Single Memory Multi Processor, shared memory multi processing unit like a multicore.
- Node : Component which are connected through the network in a distributed memory multiprocessors computers.

There is two option for getting the power on distributed computers :

- Many low power nodes : Massively Parallel Processor (MPP) → "army of ants".
- Few high power nodes : connecting SMP nodes → "herd of elephants".

Chapitre 2

Architecture

2.1 Parallel architecture

Two main class :

- Shared memory : several processors sharing through a bus or a dynamic network the same set of memory bank.
- Distributed memory : several processors, having each its own local memory, communicate through a static or dynamic network.

2.2 Shared-address-space plaforms

- UMA : Uniform Memory Access → shared-address-space computer with local caches and global memories. All memory access times (except cache) are identical.
- NUMA : Non Uniform Memory Access → shared-address-space computer with local memory only. Local memory access times are shorter.

2.3 Implicit and Explicit parallelism

Implicit parallelism :

- Higher level of device integration have made available a large number of transistors.
- Current processors use these resources in multiple functional units and execute multiple instructions in the same cycle.
- The precise manner in which these instructions are selected and executed provides impressive diversity in architectures.

Explicit parallelism :

- An explicit parallel program must specify concurrency and interaction between concurrent subtasks.

SISD MISD
SIMD MIMD

2.4 Parallel vs. Distributed computation

Parallel processing refers to multiple CPUs within the same shared-memory machine performing computation.

Distributed computation involves multiple computers with their own memory communicating over a network.

2.5 Interconnection network

static : network consist of point-to-point communication links among processing nodes.

dynamic : networks are build using switches and communication links.

2.6 Structural approach

Based on the four main components of a computer :

- The control unit
- The computing unit
- The program memory
- The data memory

2.6.1 Sequential computer

There is only one of each component : One program using one set of data. At each time step only one instruction of the program is executed by the unique processing unit.

2.6.2 Vectorial computer

There is one control unit, one program memory, multiple data memories and multiple computing units : One program executes several times the same instruction on several data (vector) in a synchronous way (one control unit).

2.6.3 Parallel computer

One data memory, several program memories, several computing units, several control units and several different programs execute asynchronously on the same data set. All programs share the same memory.

2.6.4 Distributed computer

All component are duplicated, a network is added to communicate between the different processor. Several programs execute asynchronously on different data set : Each program has its own memory and the communications between the programs are done through the network.

2.7 Network topology

2.7.1 Bus

- Some of the simplest and earliest parallel machines used buses.
- All processors access a common bus for exchanging data.
- The distance between any two nodes is $O(1)$ in a bus.
- The bus provides a convenient broadcast media.
- The bandwidth of the shared bus is a major bottleneck.
- Typical bus based machines are limited to dozens of nodes.

2.7.2 Crossbars

It's like a matrix of connexion between the node, it use a $p \times m$ grid of switches to connect p inputs to m outputs in a non-blocking manner. The cost of a crossbar of p processors grows as $O(p^2)$, so it's difficult to scale for large values of p .


2.7.3 Multistage Networks

- Buses have excellent cost scalability, but poor performance scalability.
- Crossbars gave excellent performance scalability but poor cost scalability.
- Multistage interconnects strike a compromise between these extremes.
- It consists of $\log p$ stages, where p is the number of inputs/outputs.
- At each stage, input i is connected to output j if :

$$j = \begin{cases} 2i & 0 \leq i \leq \frac{p}{2} - 1 \\ 2i + 1 - p & \frac{p}{2} \leq i \leq p - 1 \end{cases}$$


2.7.4 Multistage Omega Network

The perfect shuffle patterns are connected using 2×2 switches, there is two modes : cross-over or pass-through.




2.7.5 Stars

Every node is connected only to a common node at the center, the distance between any pair of nodes is $O(1)$. The bottleneck is the central node.




2.7.6 Linear arrays, Meshes and Torus

- In a linear array, each node has two neighbors, one to its left and one to its right.
- If the nodes at either end are connected, we refer to it as 1 – D torus or a ring.
- A generalization to $2d$ has nodes with 4 neighbors : north, south, east, west.
- A further generalization to d dimensions has nodes with $2d$ neighbors.



2.7.7 Hypercubes

- A special case of a d -dimensional mesh is a hypercube.
- $d = \log p$, where p is the total number of nodes.
- the distance between any two nodes is at most $\log p$.
- each node has $\log p$ neighbors.
- The distance between two nodes is given by the number of bit positions at which the two nodes differ.



2.7.8 K-ring

Intuitively, the K-ring topology is a graph built using K rings where each ring goes over all the nodes in a different order. The value K is called the dimension. More formally, the definition of the K -ring topology is the following :

- The size N is an integer > 0 .
- The dimension K is an integer > 0 .
- K different positive integers (a_1, \dots, a_K) , prime with N and smaller than $\frac{N+1}{2}$.
- The nodes are numbered from 0 to $N - 1$.
- Each node i is connected to the nodes $i + a_j \pmod N, j \in \{1, \dots, K\}$

Conventions :

- With such a construction each a_j defines a ring on the N nodes, this ring will be identified as the $ring_j$ and a_j will be called the $step_j$. By renumbering the nodes it is always possible, for any K -rings to obtain a $ring_j$ with a $step_j$ equal to 1. By convention, we decide that this ring is numbered by 1. Consequently we will only consider K -rings with $a_1 = 1$.

2.8 Evaluating interconnection networks

- Diameter : distance between the farthest two nodes in the network.
- Bisection width : minimum number of wires you must cut to divide the network into two equal parts.
- Cost
 - number of links or switches is a meaningful measure of the cost.
 - a number of other factors, such as the ability to layout the network, the length of wires,...
- Scalability
 - the ability to evolve.
 - the ability to be realized for any number of nodes.

Network	Diameter	Bisection	Degree	Cost
Completely-connected	1	$\frac{p^2}{4}$	$p - 1$	$\frac{p(p-1)}{2}$
Star	2	1	1	$p - 1$
Binary Tree	$2 \log(\frac{p+1}{2})$	1	1	$p - 1$
2-D Mesh	$p - 1$	1	1	$p - 1$
2-D Torus	$2\sqrt{p} - 1$	\sqrt{p}	2	$2(p - \sqrt{p})$
Hypercube	$\log p$	$\frac{p}{2}$	$\log p$	$\frac{(p \log p)}{2}$
K-D Torus	$d \lfloor \frac{k}{2} \rfloor$	$2k^{d-1}$	$2d$	dp
Crossbar	1	p	1	p^2
Omega Network	$\log p$	$\frac{p}{2}$	2	$\frac{p}{2}$
Dynamic Tree	$2 \log p$	1	2	$p - 1$

2.9 Communication cost

Overhead in parallel programs comes from :

- idling
- contention
- communication

Communication cost depends on :

- communication model
- network topology
- data handling and routing
- associated software protocols
- ...

The total time to transfer a message over a network comprises of the following :

- Startup time t_s : time spent at sending and receiving nodes (executing the routing algorithm, programming routers, etc...).
- Per-hop time t_h : function of number of hops and includes factors such as switch latencies, network delays, etc...
- Per-word time t_w : time which includes all overhaeds that are determined by the length of the message. This includes bandwidth of links, error checking and correction, etc...

The cost of communicating a message between two nodes / hops away using cut-through routing is given by $t_{comm} = t_s + lt_h + t_w m$. By the ways, lt_h is smaller than t_s and t_w so we generally have $t_{comm} = t_s + t_w m$.

For Shared Address Space Machines, the simplified cost model is still valid, but a number of other factors make accurate cost modeling more difficult :

- memory layout determined by the system
- finite cache size can result in cache thrashing
- overheads associated with invalidate and update operations are difficult to quantify
- spatial locality is difficult to model
- prefetching can play a role in reducing the overhead associated with data access
- false sharing and contention are difficult to model

Chapitre 3

Models

Chapitre 4

Communication